

ANLP - Assignment - 1

Brief explanation of solution:

Sentiment analysis is the process of determining the emotional tone of a piece of text. One common solution for sentiment analysis is to use machine learning algorithms to classify text as positive, negative, or neutral.

In this assignment I used perceptron, SVM, Logistic Regression, and Naive Bayes algorithms. I first preprocessed the text data by cleaning and later tokenizing it, and then used techniques such as TF-IDF to convert the text into numerical feature representations that can be input into a machine learning model.

As part of data cleaning, removing irrelevant, incorrect, or inconsistent data from the dataset using the following methods: Reviews don't follow the conventional grammatical structure, so require the following cleaning to be done.

1. Converted all strings to lowercase to interpret both lowercase and capitalized words as the same
2. Removed URLs as the url might contain new information but the url text does not add any value to the text - Used regex to match urls
3. Removed HTML tags as they do not aid in classification - used BeautifulSoup html parser
4. Removed contractions for consistency and simplicity - used contractions library
5. Removed non-alphabetic characters and numbers as they can add noise to the feature set and make the model less accurate - Used regex
6. Removed extra spaces to save on memory and storage space - Used regex
7. Converted emojis to text to better represent the sentiment - Used EMOTICONS_EMO library

Before performing sentiment analysis, it is crucial to preprocess the text data.

Below methods are used for data preprocessing:

1. Text cleaning: cleaning the text by removing stop words (nltk.corpus.stopwords) and lemmatization (using WordNetLemmatizer).
Created two sets of reviews - one by removing stopwords (dataframe - df) and other including stopwords (dataframe - df2)
2. Tokenization: It is performed using TfidfVectorizer by breaking down the text into individual words or phrases to better understand the context and meaning of the words. Used n-grams (specifically 3-gram) to group words to get context.
3. Text encoding: Text encoding is a process of converting the text into numerical representations that can be input into a machine learning model. TfidfVectorizer also does this.

Later sampled data into train and test set with 80-20% split. Split is done such that the train and test contain an equal number of samples for each class.

Once the text is preprocessed, the model is trained on a labeled dataset of text and corresponding sentiments. Cross-validation techniques are used to evaluate the model's performance and the trained model is then used to classify new, unseen text data.

Used LinearSVC for SVM classifier (simple and faster to train), and MultinomialNB (it is an alternative to the "heavy" AI-based semantic analysis and drastically simplifies textual data classification) for naive bayes.

Preprocessing the text data before sentiment analysis ensures that the machine learning model is trained on clean, consistent and useful data which leads to better performance and accuracy.

Got below results:

Stopwords are removed in preprocessing for computing average length:

- 1) Average length of reviews in terms of character length **before and after** data cleaning:
- 2) Average length of reviews in terms of character length **before and after** Preprocessing:

Metrics with removing stopwords in preprocessing:

Perceptron:

Class 1: Precision - 0.673538740371545 , Recall - 0.74325 , F1-score - 0.7066793439505585
Class 2: Precision - 0.6193907631837536 , Recall - 0.47275 , F1-score - 0.5362257195519636
Class 3: Precision - 0.7182881094198103 , Recall - 0.814 , F1-score - 0.7631548107347942
Average: Precision - 0.6766666666666666 , Recall - 0.6766666666666666 , F1-score - 0.6766666666666666

SVM:

Class 1: Precision - 0.7066477407144547 , Recall - 0.74675 , F1-score - 0.7261456180867875
Class 2: Precision - 0.6310975609756098 , Recall - 0.56925 , F1-score - 0.5985804416403786
Class 3: Precision - 0.7771908763505402 , Recall - 0.80925 , F1-score - 0.792896509491733
Average: Precision - 0.7084166666666667 , Recall - 0.7084166666666667 , F1-score - 0.7084166666666667

Logistic Regression:

Class 1: Precision - 0.7166585246702492 , Recall - 0.7335 , F1-score - 0.7249814677538918
Class 2: Precision - 0.6244518957957184 , Recall - 0.60525 , F1-score - 0.6147010283102704
Class 3: Precision - 0.7803425167535368 , Recall - 0.786 , F1-score - 0.7831610412255572
Average: Precision - 0.70825 , Recall - 0.70825 , F1-score - 0.70825

Naive Bayes:

Class 1: Precision - 0.7185148018063221 , Recall - 0.716 , F1-score - 0.7172551965940395

Class 2: Precision - 0.6042411246128186 , Recall - 0.634 , F1-score - 0.61876296205929
Class 3: Precision - 0.7982708933717579 , Recall - 0.76175 , F1-score - 0.7795829602149162
Average: Precision - 0.7039166666666666 , Recall - 0.7039166666666666 , F1-score - 0.7039166666666666

Metrics without removing stopwords in preprocessing:

Perceptron:

Class 1: Precision, Recall, F1-score - 0.7090085795996187, 0.74375, 0.7259638848218644
Class 2: Precision, Recall, F1-score - 0.648235294117647, 0.551, 0.5956756756756757
Class 3: Precision, Recall, F1-score - 0.7606721162579473, 0.8375, 0.7972394098048549
Average: Precision, Recall, F1-score - 0.71075, 0.71075, 0.71075

SVM:

Class 1: Precision, Recall, F1-score - 0.739343459088682 0.7545 0.746844840386043
Class 2: Precision, Recall, F1-score - 0.6593085106382979 0.61975 0.6389175257731958
Class 3: Precision, Recall, F1-score - 0.8027898027898028 0.8345 0.8183378278989948
Average: Precision, Recall, F1-score - 0.73625 0.73625 0.73625

Logistic Regression:

Class 1: Precision, Recall, F1-score - 0.7415338645418327 0.7445 0.7430139720558881
Class 2: Precision, Recall, F1-score - 0.6545500762582613 0.64375 0.6491051172170407
Class 3: Precision, Recall, F1-score - 0.8037037037037037 0.81375 0.8086956521739131
Average: Precision, Recall, F1-score - 0.734 0.734 0.734

Naive Bayes:

Class 1: Precision, Recall, F1-score - 0.7529777317452098 0.727 0.7397608750953956
Class 2: Precision, Recall, F1-score - 0.6170583115752829 0.709 0.6598417868776175
Class 3: Precision, Recall, F1-score - 0.8565782044042913 0.7585 0.8045611243701936
Average: Precision, Recall, F1-score - 0.7315 0.7315 0.7315

Observation:

Got better metrics by including stopwords - obviously including would convey the meaning but it is computational expensive and redundant.

Out of the four models, perceptron metrics are the lowest. SVM, Logistic regression, and Naive Bayes got similar results but among these SVM got the highest score for precision.

Import libraries

```
In [70]: import pandas as pd
import numpy as np
import nltk
nltk.download('wordnet')
import re
from bs4 import BeautifulSoup
import os
import contractions
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('omw-1.4')

from sklearn.metrics import precision_score, recall_score, f1_score, classification_report

import re
import pickle
from emot.emo_unicode import UNICODE_EMOJI # For emojis
from emot.emo_unicode import EMOTICONS_EMO # For EMOTICONS

[nltk_data] Downloading package wordnet to
[nltk_data]   /Users/venkatasaisumanthsadu/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data]   /Users/venkatasaisumanthsadu/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   /Users/venkatasaisumanthsadu/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package omw-1.4 to
[nltk_data]   /Users/venkatasaisumanthsadu/nltk_data...
[nltk_data]   Package omw-1.4 is already up-to-date!
```

```
In [71]: !pip install bs4 # in case you don't have it installed
!pip install emot
```

Requirement already satisfied: bs4 in /Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages (0.0.1)
Requirement already satisfied: beautifulsoup4 in /Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages (from bs4) (4.11.1)
Requirement already satisfied: soupsieve>1.2 in /Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages (from beautifulsoup4->bs4) (2.3.1)
Requirement already satisfied: emot in /Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages (3.1)

Read Data

```
In [72]: #to get the current working directory
directory = os.getcwd()
url = os.path.join(directory, "data.tsv")
df = pd.read_csv(url, sep='\t', header=0, on_bad_lines='skip')

/var/folders/dz/k6x51lvd2jv050r966v_h2br0000gn/T/ipykernel_6120/1309074324.py:4: DtypeWarning: Columns (8) have mixed
types. Specify dtype option on import or set low_memory=False.
    df = pd.read_csv(url, sep='\t', header=0, on_bad_lines='skip')
```

Keep Reviews and Ratings

```
In [73]: df = df[['review_body', 'star_rating']]
```

We form three classes and select 20000 reviews randomly from each class.

```
In [74]: df = df[(df['star_rating'].eq(1) | df['star_rating'].eq(2) | df['star_rating'].eq(3) | df['star_rating'].eq(4) | df['star_rating'].eq(5)) & (df['class'] = df['star_rating'].apply(lambda x: 1 if x in [1, 2] else 2 if x == 3 else 3))
df.head(2)
```

Out[74]:

	review_body	star_rating	class
32768	I have a bunch of these color-changing gel pol...	5	3
32769	this product smells like it's been dipped in f...	1	1

```
In [75]: # there are few nan values in the dataframe, also removing duplicate rows
df = df.dropna()
df = df.drop_duplicates()
df = df.reset_index()
```

```
In [76]: class1 = df[df['class']==1].sample(n=20000, random_state=42)
class2 = df[df['class']==2].sample(n=20000, random_state=42)
class3 = df[df['class']==3].sample(n=20000, random_state=42)
df = pd.concat([class1, class2, class3])
```

```
In [77]: # average length before data cleaning:
print('Average length of the reviews before data cleaning :', (df['review_body'].str.len()).mean())
```

Average length of the reviews before data cleaning : 289.2713

Data Cleaning

```
In [78]: def remove_urls(text):
    text = re.sub(r'(\https|http)?:\:\/\/(\w|\.|\/|\?|\=|&|\%)*\b', '', text, flags=re.MULTILINE)
    return (text)

def remove_contractions(text) :
    expanded_words = []
    for word in text.split():
        expanded_words.append(contractions.fix(word))
    expanded_text = ' '.join(expanded_words)
    return expanded_text

remove_non_english = lambda s: re.sub(r'^a-zA-Z', ' ', s)
remove_spaces = lambda s: re.sub(' +', ' ', s)
```

```
In [79]: def cleaning(text):
    #remove urls
    text = remove_urls(text)
    #remove html tags
    text = BeautifulSoup(text, "lxml").text
    #remove contractions
    text = remove_contractions (text)
    #remove non-alphabetic chars
    text = remove_non_english(text)
    #lowercase
    text = text.lower( )
    #remove extra spaces
    text = remove_spaces(text)

    return text
```

```
In [80]: df['cleaned_text_reviews'] = list(map(cleaning, df.review_body))
```

/Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages/bs4/__init__.py:435: MarkupResemblesLocatorWarning: The input looks more like a filename than markup. You may want to open this file and pass the filehandle into BeautifulSoup.
warnings.warn(

```
In [81]: def convert_emojis(text):
    for emot in UNICODE_EMOJI:
        text = text.replace(emot, "_".join(UNICODE_EMOJI[emot].replace(",","").replace(":","").split()))
    return text

df['cleaned_text_reviews'] = df['cleaned_text_reviews'].apply(lambda row: convert_emojis(str(row)))
```

```
In [82]: # average length after data cleaning:
print('Average length of the reviews after data cleaning :', (df['cleaned_text_reviews'].str.len()).mean())
```

Average length of the reviews after data cleaning : 279.3924166666667

Pre-processing

```
In [83]: # average length before pre-processing:
print('Average length of the reviews before pre-processing :', (df['cleaned_text_reviews'].str.len()).mean())
```

Average length of the reviews before pre-processing : 279.3924166666667

```
In [84]: df2 = df.copy()
```

remove the stop words

```
In [85]: from nltk.corpus import stopwords

to_remove = ['not']
new_stopwords = set(stopwords.words('english')).difference(to_remove)

df2['cleaned_text_reviews'] = df2['cleaned_text_reviews'].apply(lambda x: " ".join(x for x in x.split() if x not in new_stopwords))
```

perform lemmatization

```
In [86]: from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()

df2['cleaned_text_reviews'] = df2['cleaned_text_reviews'].apply(lambda x: " ".join([lemmatizer.lemmatize(word) for word in x.split()]))
```

```
In [87]: # average length after pre-processing:
print('Average length of the reviews after pre-processing :', (df2['cleaned_text_reviews'].str.len()).mean())

Average length of the reviews after pre-processing : 172.00831666666667
```

TF-IDF Feature Extraction

```
In [88]: from sklearn.model_selection import train_test_split
x_train, x_valid, y_t, y_v = train_test_split(df2['cleaned_text_reviews'], df2['class'], test_size=0.2, stratify=df2['class'])
```

```
In [89]: from sklearn.feature_extraction.text import TfidfVectorizer

tfidf = TfidfVectorizer(ngram_range=(1, 3))
tfidf.fit(df2['cleaned_text_reviews'])
```

```
Out[89]: TfidfVectorizer(ngram_range=(1, 3))
```

```
In [90]: x_t = tfidf.transform(x_train)
x_v = tfidf.transform(x_valid)
```

Perceptron

```
In [91]: from sklearn.linear_model import Perceptron
p = Perceptron()
p.fit(x_t, y_t)
```

```
Out[91]: Perceptron()
```

```
In [92]: report = classification_report(y_v, p.predict(x_v), output_dict=True)
```

```
In [93]: print('Perceptron:')
print('Class 1: Precision - ', report['1']['precision'], ', Recall - ', report['1']['recall'], ', F1-score - ', report['1']['f1_score'], '\n')
print('Class 2: Precision - ', report['2']['precision'], ', Recall - ', report['2']['recall'], ', F1-score - ', report['2']['f1_score'], '\n')
print('Class 3: Precision - ', report['3']['precision'], ', Recall - ', report['3']['recall'], ', F1-score - ', report['3']['f1_score'], '\n')
print('Average: Precision - ', precision_score(y_v, p.predict(x_v), average='micro'), ', Recall - ', recall_score(y_v, p.predict(x_v), average='micro'), ', F1-score - ', f1_score(y_v, p.predict(x_v), average='micro'))
```

```
Perceptron:
Class 1: Precision - 0.673538740371545, Recall - 0.74325, F1-score - 0.7066793439505585
Class 2: Precision - 0.6193907631837536, Recall - 0.47275, F1-score - 0.5362257195519636
Class 3: Precision - 0.7182881094198103, Recall - 0.814, F1-score - 0.7631548107347942
Average: Precision - 0.6766666666666666, Recall - 0.6766666666666666, F1-score - 0.6766666666666666
```

SVM

```
In [94]: from sklearn.svm import LinearSVC
sv = LinearSVC()
sv.fit(x_t, y_t)
```

```
Out[94]: LinearSVC()
```

```
In [95]: report_sv = classification_report(y_v, sv.predict(x_v), output_dict=True )
```

```
In [96]: print('SVM:')
print('Class 1: Precision - ', report_sv['1']['precision'], ', Recall - ', report_sv['1']['recall'], ', F1-score - ', r
print('Class 2: Precision - ', report_sv['2']['precision'], ', Recall - ', report_sv['2']['recall'], ', F1-score - ', r
print('Class 3: Precision - ', report_sv['3']['precision'], ', Recall - ', report_sv['3']['recall'], ', F1-score - ', r
print('Average: Precision - ', precision_score(y_v, sv.predict(x_v), average='micro'), ', Recall - ', recall_score(y_v,
```

```
SVM:
Class 1: Precision - 0.7066477407144547 , Recall - 0.74675 , F1-score - 0.7261456180867875
Class 2: Precision - 0.6310975609756098 , Recall - 0.56925 , F1-score - 0.5985804416403786
Class 3: Precision - 0.7771908763505402 , Recall - 0.80925 , F1-score - 0.792896509491733
Average: Precision - 0.7084166666666667 , Recall - 0.7084166666666667 , F1-score - 0.7084166666666667
```

Logistic Regression

```
In [97]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_t, y_t)
```

/Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[97]: LogisticRegression()
```

```
In [98]: report_lr = classification_report(y_v, lr.predict(x_v), output_dict=True )
```

```
In [99]: print('Logistic Regression:')
print('Class 1: Precision - ', report_lr['1']['precision'], ', Recall - ', report_lr['1']['recall'], ', F1-score - ', r
print('Class 2: Precision - ', report_lr['2']['precision'], ', Recall - ', report_lr['2']['recall'], ', F1-score - ', r
print('Class 3: Precision - ', report_lr['3']['precision'], ', Recall - ', report_lr['3']['recall'], ', F1-score - ', r
print('Average: Precision - ', precision_score(y_v, lr.predict(x_v), average='micro'), ', Recall - ', recall_score(y_v,
```

```
Logistic Regression:
Class 1: Precision - 0.7166585246702492 , Recall - 0.7335 , F1-score - 0.7249814677538918
Class 2: Precision - 0.6244518957957184 , Recall - 0.60525 , F1-score - 0.6147010283102704
Class 3: Precision - 0.7803425167535368 , Recall - 0.786 , F1-score - 0.7831610412255572
Average: Precision - 0.70825 , Recall - 0.70825 , F1-score - 0.70825
```

Naive Bayes

```
In [100]: from sklearn.naive_bayes import MultinomialNB
nb = MultinomialNB()
nb.fit(x_t, y_t)
```

```
Out[100]: MultinomialNB()
```

```
In [101]: report_nb = classification_report(y_v, nb.predict(x_v), output_dict=True )
```

```
In [102]: print('Naive Bayes:')
print('Class 1: Precision - ', report_nb['1']['precision'], ', Recall - ', report_nb['1']['recall'], ', F1-score - ', r
print('Class 2: Precision - ', report_nb['2']['precision'], ', Recall - ', report_nb['2']['recall'], ', F1-score - ', r
print('Class 3: Precision - ', report_nb['3']['precision'], ', Recall - ', report_nb['3']['recall'], ', F1-score - ', r
print('Average: Precision - ', precision_score(y_v, nb.predict(x_v), average='micro'), ', Recall - ', recall_score(y_v,
```

```
Naive Bayes:
Class 1: Precision - 0.7185148018063221 , Recall - 0.716 , F1-score - 0.7172551965940395
Class 2: Precision - 0.6042411246128186 , Recall - 0.634 , F1-score - 0.61876296205929
Class 3: Precision - 0.7982708933717579 , Recall - 0.76175 , F1-score - 0.7795829602149162
Average: Precision - 0.7039166666666666 , Recall - 0.7039166666666666 , F1-score - 0.7039166666666666
```

```
In [ ]:
```

In []:

Model training without doing stopwords preprocessing

perform lemmatization

```
In [103]: from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()

df['cleaned_text_reviews'] = df['cleaned_text_reviews'].apply(lambda x: " ".join([lemmatizer.lemmatize(word) for word in
```

TF-IDF Feature Extraction

```
In [104]: from sklearn.model_selection import train_test_split
x_train2 ,x_valid2 ,y_t2 ,y_v2 = train_test_split(df['cleaned_text_reviews'],df['class'],test_size=0.2, stratify = df['

from sklearn.feature_extraction.text import TfidfVectorizer

tfidf = TfidfVectorizer(ngram_range=(1, 3))
tfidf.fit(df['cleaned_text_reviews'])
x_t2 = tfidf.transform(x_train2)
x_v2 = tfidf.transform(x_valid2)
```

Perceptron

```
In [105]: from sklearn.linear_model import Perceptron
p = Perceptron()
p.fit(x_t2, y_t2)
report = classification_report(y_v2, p.predict(x_v2), output_dict=True )
```

```
In [106]: print('Perceptron')
print('Class 1: Precision, Recall, F1-score - ', report['1']['precision'],report['1']['recall'],report['1']['f1-score'])
print('Class 2: Precision, Recall, F1-score - ', report['2']['precision'],report['2']['recall'],report['2']['f1-score'])
print('Class 3: Precision, Recall, F1-score - ', report['3']['precision'],report['3']['recall'],report['3']['f1-score'])
print('Average: Precision, Recall, F1-score - ', precision_score(y_v2, p.predict(x_v2), average='micro'), recall_score
```

```
Perceptron
Class 1: Precision, Recall, F1-score - 0.7090085795996187 0.74375 0.7259638848218644
Class 2: Precision, Recall, F1-score - 0.648235294117647 0.551 0.5956756756756757
Class 3: Precision, Recall, F1-score - 0.7606721162579473 0.8375 0.7972394098048549
Average: Precision, Recall, F1-score - 0.71075 0.71075 0.71075
```

SVM

```
In [107]: from sklearn.svm import LinearSVC
sv = LinearSVC()
sv.fit(x_t2, y_t2)
report_sv = classification_report(y_v2, sv.predict(x_v2), output_dict=True )
```

```
In [108]: print('SVM:')
print('Class 1: Precision, Recall, F1-score - ', report_sv['1']['precision'],report_sv['1']['recall'],report_sv['1']['f
print('Class 2: Precision, Recall, F1-score - ', report_sv['2']['precision'],report_sv['2']['recall'],report_sv['2']['f
print('Class 3: Precision, Recall, F1-score - ', report_sv['3']['precision'],report_sv['3']['recall'],report_sv['3']['f
print('Average: Precision, Recall, F1-score - ', precision_score(y_v2, sv.predict(x_v2), average='micro'), recall_score
```

```
SVM:
Class 1: Precision, Recall, F1-score - 0.739343459088682 0.7545 0.746844840386043
Class 2: Precision, Recall, F1-score - 0.6593085106382979 0.61975 0.6389175257731958
Class 3: Precision, Recall, F1-score - 0.8027898027898028 0.8345 0.8183378278989948
Average: Precision, Recall, F1-score - 0.73625 0.73625 0.73625
```


Logistic Regression

```
In [109]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_t2, y_t2)
report_lr = classification_report(y_v2, lr.predict(x_v2), output_dict=True )

/Users/venkatasaisumanthsadu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:814: Converge
enceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(

In [110]: print('Logistic Regression:')
print('Class 1: Precision, Recall, F1-score - ', report_lr['1']['precision'],report_lr['1']['recall'],report_lr['1']['f
print('Class 2: Precision, Recall, F1-score - ', report_lr['2']['precision'],report_lr['2']['recall'],report_lr['2']['f
print('Class 3: Precision, Recall, F1-score - ', report_lr['3']['precision'],report_lr['3']['recall'],report_lr['3']['f
print('Average: Precision, Recall, F1-score - ', precision_score(y_v2, lr.predict(x_v2), average='micro'), recall_score

Logistic Regression:
Class 1: Precision, Recall, F1-score - 0.7415338645418327 0.7445 0.7430139720558881
Class 2: Precision, Recall, F1-score - 0.6545500762582613 0.64375 0.6491051172170407
Class 3: Precision, Recall, F1-score - 0.8037037037037037 0.81375 0.8086956521739131
Average: Precision, Recall, F1-score - 0.734 0.734 0.734
```

Naive Bayes

```
In [111]: from sklearn.naive_bayes import MultinomialNB
nb = MultinomialNB()
nb.fit(x_t2, y_t2)
report_nb = classification_report(y_v2, nb.predict(x_v2), output_dict=True )

In [112]: print('Naive Bayes:')
print('Class 1: Precision, Recall, F1-score - ', report_nb['1']['precision'],report_nb['1']['recall'],report_nb['1']['f
print('Class 2: Precision, Recall, F1-score - ', report_nb['2']['precision'],report_nb['2']['recall'],report_nb['2']['f
print('Class 3: Precision, Recall, F1-score - ', report_nb['3']['precision'],report_nb['3']['recall'],report_nb['3']['f
print('Average: Precision, Recall, F1-score - ', precision_score(y_v2, nb.predict(x_v2), average='micro'), recall_score

Naive Bayes:
Class 1: Precision, Recall, F1-score - 0.7529777317452098 0.727 0.7397608750953956
Class 2: Precision, Recall, F1-score - 0.6170583115752829 0.709 0.6598417868776175
Class 3: Precision, Recall, F1-score - 0.8565782044042913 0.7585 0.8045611243701936
Average: Precision, Recall, F1-score - 0.7315 0.7315 0.7315
```